

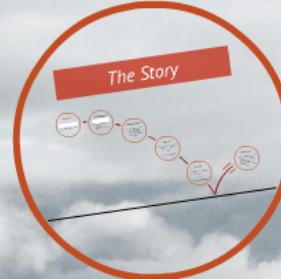
# Lucene Server

*From Erlang to Java and Back Again*



# Lucene Server

*From Erlang to Java and Back Again*





# Hello World!!

## About Me

- Programmer since the age of 10
- Coded in Basic, .NET, Java, JS, ...
- Erlang programmer for the last 5 years
- Director of Engineering at Inaka

## About Inaka

- We're based in Argentina
- We do Erlang consulting
- We build end-to-end applications
- We develop highly-concurrent servers for applications like Whisper or TigerText

## About TigerText

- Provides secure communication for mobile devices
- Heavily used by health-care companies
- Also provides the underlying communication platform for other apps
- All servers written in Erlang



## *About Me*

- Programmer since the age of 10
- Coded in Basic, .NET, Java, JS, ...
- Erlang programmer for the last 5 years
- Director of Engineering at Inaka



## *About Inaka*

- We're based in Argentina
- We do Erlang consulting
- We build end-to-end applications
- We develop highly-concurrent servers  
for applications like Whisper or  
TigerText



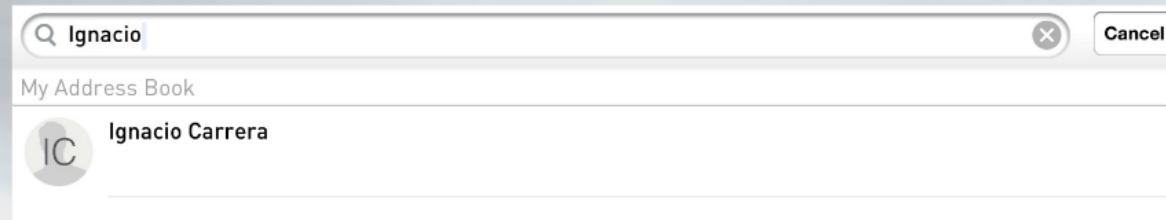
## *About TigerText*

- Provides secure communication for mobile devices
- Heavily used by health-care companies
- Also provides the underlying communication platform for other apps
- All servers written in Erlang

# *The Story*

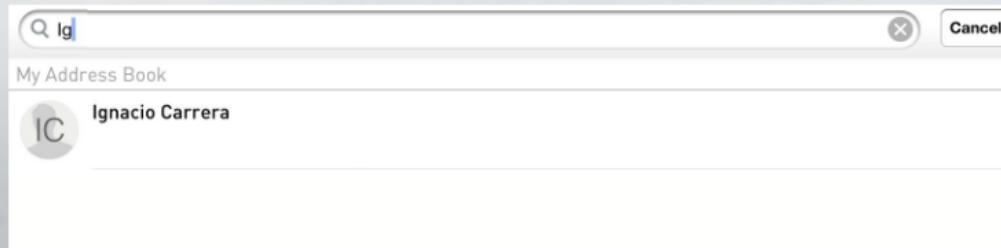


# The Scenario



We had a very basic search field for users...

# Functional Requests



We wanted to add the following features:

- Partial Matching
- Autocomplete
- Advanced Search (by field, location, etc.)
- Ranked Results

# Technical Requests

We also wanted our solution to be:

- Implemented in Erlang
- Included in our existing servers as a rebar dependency
- Fast and scalable

# **Lucene Core**

A high-performance, scalable, full-featured text search engine library with powerful, accurate and efficient search algorithms

...but...

**It's written in JAVA**

# JInterface

A Java library that let's you create "nodes" in Java  
and provides the means to communicate them  
with regular Erlang nodes

...and...

It comes with Erlang!

# **Lucene Server**

An OTP compliant application  
that provides indexing and  
querying access to a lucene  
backend without any JAVA  
knowledge requirement for its  
users



# *Lucene Server*

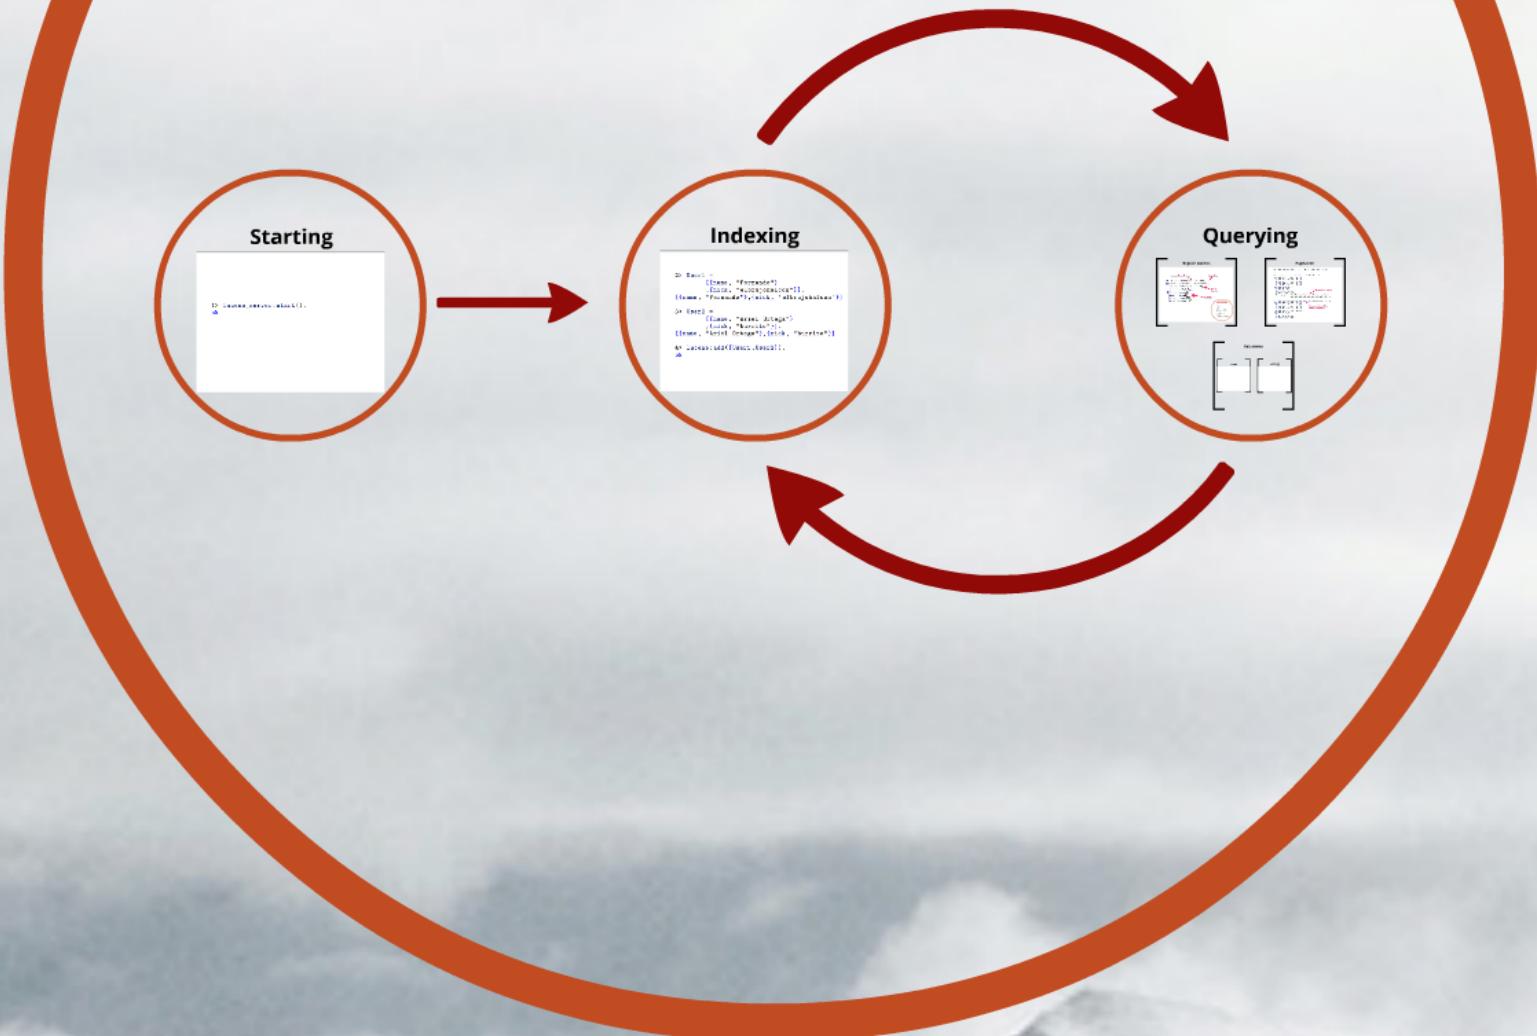
How it works



How it's done



# How it works



# Starting

```
1> lucene_server:start().  
ok
```

# Indexing

```
2> User1 =  
    [{name, "Fernando"}  
     ,{nick, "elbrujohalcon"}].  
[{name, "Fernando"},{nick, "elbrujohalcon"}]  
  
3> User2 =  
    [{name, "Ariel Ortega"}  
     ,{nick, "burrito"}].  
[{name, "Ariel Ortega"},{nick, "burrito"}]  
  
4> lucene:add([User1,User2]).  
ok
```

# Querying

## Regular Queries

```
Lucene Query Syntax  
5> lucene:match("nick: b*", 10).  
[[{"name": "Fernando Benavides",  
    "nick": "elbrujohalcon",  
    "score": 1.0}],  
[{"total_hits": 1},  
 {"first_hit": 1},  
 {"query_time": 7837},  
 {"search_time": 457}]]
```



## Pagination

```
5> lucene:and([{{code, X}} || X <- lists:seq(1,10)]).  
ok  
#> [_, M] = lucene:match("code:[0 TO 10]", 5).  
[[{code:1}, {"score":1.0}],  
 [{code:2}, {"score":1.0}],  
 [{code:3}, {"score":1.0}],  
 [{code:4}, {"score":1.0}],  
 [{code:5}, {"score":1.0}],  
 [{"total_hits":10},  
 {"first_hit":1},  
 {"query_time":13539},  
 {"search_time":28115},  
 {"last_page":<>[172,237,0,5,115,114,0,36,99,111,109,46,  
 116,108,103,101,114,116,101,120,...>]}]  
  
5> lucene:continuer(groupsize:get_value(next_page, M), 5).  
[[{code:4}, {"score":1.0}],  
 [{code:5}, {"score":1.0}],  
 [{code:6}, {"score":1.0}],  
 [{code:7}, {"score":1.0}],  
 [{"total_hits":10},  
 {"first_hit":6},  
 {"query_time":23600},  
 {"search_time":17311}]]
```

Token for Next Page  
It's used in the call to  
lucene:continuer/2

## Extensions

.near

```
5> lucene:match("nick:near(ben)", 10).  
[[{"name": "Fernando Benavides",  
    "nick": "elbrujohalcon",  
    "score": 1.0}],  
[{"total_hits": 1},  
 {"first_hit": 1},  
 {"query_time": 7837},  
 {"search_time": 457}]]
```

.erlang

```
5> lucene:match("nick:erlang", 10).  
[[{"name": "Fernando Benavides",  
    "nick": "elbrujohalcon",  
    "score": 1.0}],  
[{"total_hits": 1},  
 {"first_hit": 1},  
 {"query_time": 7837},  
 {"search_time": 457}]]
```

# Regular Queries

```
Lucene Query Syntax      Page Size  
5> lucene:match("nick: b*", 10).  
{[[[ {name , "Fernando Benavides"} ,  
      {nick , "elbrujohalcon"} ,  
      { 'score' , 1.0 } ]],  
   [{total_hits , 1} ,  
    {first_hit , 1} ,  
    {query_time , 783} ,  
    {search_time , 457} ]}
```

List of Results

Metadata

## Lucene Query Syntax

- Terms:
  - 'Lucene'
- Fields:
  - 'title:"Lucene" AND text:test'
- Wildcards:
  - 'title:Lucene\*' OR text:text?'
- Ranges:
  - 'ranking: [5 TO 10]'
- Boosting:
  - 'title:Lucene^4 OR title:Server^1'

## *Lucene Query Syntax*

- Terms:
  - 'Lucene'
- Fields:
  - 'title:"Lucene" AND text:test'
- Wildcards:
  - 'title:Lucene\*' OR text:te?t'
- Ranges:
  - 'ranking: [5 TO 10]'
- Boosting:
  - 'title:Lucene^4 OR title:Server^1'

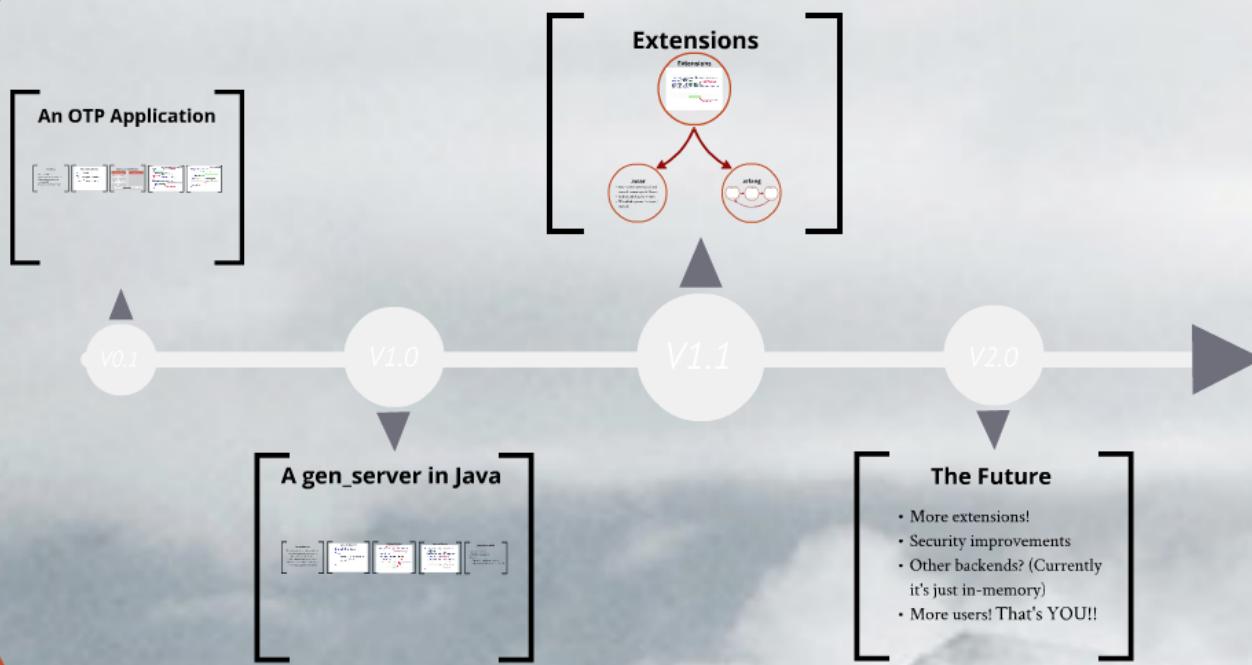
# Pagination

```
6> lucene:add([[{{code, X}]] || X <- lists:seq(1,10)]).  
ok  
  
7> {_, M} = lucene:match("code:[0 TO 10]", 5).  
{{{code,1},{score,1.0}},  
 {{code,2},{score,1.0}},  
 {{code,3},{score,1.0}},  
 {{code,4},{score,1.0}},  
 {{code,5},{score,1.0}}},  
 [{total_hits,10},  
 {first_hit,1},  
 {query_time,14335},  
 {search_time,3811},  
 {next_page,<<172,237,0,5,115,114,0,36,99,111,109,46,  
 116,105,103,101,114,116,101,120,...>>}]}  
  
8> lucene:continue(proplists:get_value(next_page, M), 5).  
{{{code,6},{score,1.0}},  
 {{code,7},{score,1.0}},  
 {{code,8},{score,1.0}},  
 {{code,9},{score,1.0}},  
 {{code,10},{score,1.0}}},  
 [{total_hits,10},  
 {first_hit,6},  
 {query_time,2368},  
 {search_time,1731}]}
```

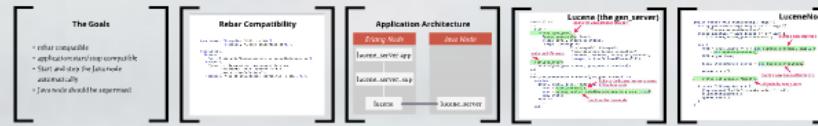
Token for Next Page

It's used in the call to  
lucene:continue/2

# How it's done



# An OTP Application



V0.1

# The Goals

- rebar compatible
- application:start/stop compatible
- Start and stop the Java node automatically
- Java node should be supervised

# Rebar Compatibility

```
{pre_hooks, [{compile, "mkdir -p bin"},  
             {compile, "./copy-jinterface.sh"}]}.  
  
{post_hooks,  
 [{clean,  
   "rm -rf bin priv/lucene_server.jar priv/OtpErlang.jar"},  
  {compile,  
   "javac -g -deprecation -sourcepath java_src  
     -classpath ./bin:/priv/* -d bin  
     java_src/**/*/*.java"},  
  {compile, "jar cf priv/lucene-server.jar -C bin ."}]}.
```

# Application Architecture

*Erlang Node*

lucene\_server app

lucene\_server\_sup

lucene

*Java Node*

lucene\_server

# Lucene (the gen\_server)

```
init([]) ->
    ...
    Port =
        erlang:open_port(
            {spawn_executable, Java},
            [{line,1000}, stderr_to_stdout,
            {args, JavaArgs ++ [{"classpath", Classpath,
                "com.tigertext.lucene.LuceneNode",
                ThisNode, JavaNode, erlang:get_cookie(),
                integer_to_list(AllowedThreads)}]}]),
    waits until it's ready
    wait_for_ready(
        #state{java_port = Port, java_node = JavaNode})
end.

...
wait_for_ready(State = #state{java_port = Port}) ->
    receive
        {Port, {data, {eol, "READY"}}, links to the lucene_server process  
in the Java node}
        true = link(process()),
        true = erlang:monitor_node(State#state.java_node, true),
        {ok, State};
        Info ->
            ...
    end.
```

**monitors the Java node**

# LuceneNode

```
public static void main(String[] args) {  
    String peerName = args.length >= 1 ? args[0]  
        : "lucene_server@localhost";  
    String nodeName = args.length >= 2 ? args[1]  
        : "lucene_server_java@localhost"; Starts a new OtpNode  
  
    try {  
        NODE = args.length >= 3 ? new OtpNode(nodeName, args[2])  
            : new OtpNode(nodeName);  
        PEER = peerName;  
  
        final OtpGenServer server = new LuceneServer(NODE);  
        server.start();  
        System.out.println("READY"); Starts a new LuceneServer in it  
    } catch (IOException e1) {  
        jlog.severe("Couldn't create node: " + e1);  
        e1.printStackTrace();  
        System.exit(1);  
    }  
}
```

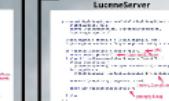
V1.0

# A gen\_server in Java

[  
Introduction  
We had a pointer on the Erlang code that received messages and, for some of them, it "returned" a value.  
That looked a lot like a gen\_server.  
We also wanted all the functionality gen\_server provides for its clients

[  
code in gen.erl  


[  
CtpGenServer  


[  
LuceneServer  


[  
JInterface-stdlib  

- Open-source Erlang
- OTF/Hotswap compatible
- Includes
  - » OptiSyncer like gen module
  - » OptiServer like gen\_server module

# Introduction

We had a process on the Java node that received messages and, for some of them, it "returned" a value.

That looked a lot like a gen\_server.

We also wanted all the functionality gen\_server provides for its clients

# code in gen.erl

```
do_call(Process, Label, Request, Timeout) ->
    try erlang:monitor(process, Process) of
        Mref ->
    ...
    catch
        error:_ ->
            %% Node (C/Java?) is not supporting the monitor.
            %% The other possible case -- this node is not
            %% distributed -- should have been handled earlier.
            %% Do the best possible with monitor_node/2.
            %% This code may hang indefinitely if the Process
            %% does not exist. It is only used for featureweak
            %% remote nodes.
            Node = get_node(Process),
            monitor_node(Node, true),
    ...
    end.
```

# OtpGenServer

```
abstract class == behavior  
public abstract class OtpGenServer extends OtpSysProcess {  
    protected OtpGenServer(OtpNode host) {  
        super(host);  
    }  
  
    protected OtpGenServer(OtpNode host, String name) {  
        super(host, name);  
    }  
  
    ...  
    protected abstract OtpErlangObject handleCall(  
        OtpErlangObject cmd, OtpErlangTuple from)  
        throws OtpStopException, OtpContinueException,  
        OtpErlangException;  
  
    protected abstract void handleCast(OtpErlangObject cmd)  
        throws OtpStopException, OtpErlangException;  
  
    protected abstract void handleInfo(OtpErlangObject cmd)  
        throws OtpStopException;  
  
    protected abstract void terminate(OtpErlangException oee);  
}
```

constructors == start\_link / init

abstract methods == behaviour callbacks

# LuceneServer

```
protected OtpErlangObject handleCall(OtpErlangObject cmd,
    OtpErlangTuple from)
throws OtpStopException, OtpContinueException,
OtpErlangException {

    OtpErlangTuple cmdTuple = (OtpErlangTuple) cmd;
    OtpErlangAtom cmdName = (OtpErlangAtom) cmdTuple.elementAt(0);

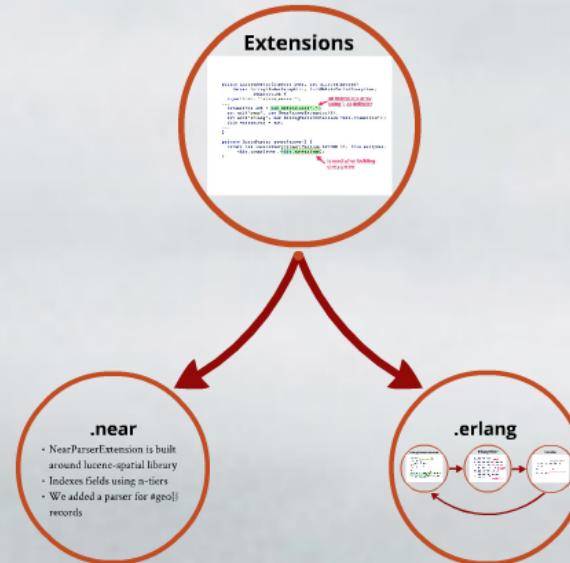
    if (cmdName.atomValue().equals("pid")) { ← {pid}
        return super.getSelf(); ← {reply, self(), State}
    } else if (cmdName.atomValue().equals("match")) {
        String queryString =
            ((OtpErlangString) cmdTuple.elementAt(1)).stringValue();
        int pageSize =
            ((OtpErlangLong) cmdTuple.elementAt(2)).intValue();

        runMatch(queryString, pageSize, from); ← {match, QueryString, PageSize}
        throw new OtpContinueException();
    } else
        ... ← {noreply, State}
}
```

# JInterface-stdlib

- Open-source library
- OTP/rebar compatible
- Includes:
  - OtpSysProcess: like gen module
  - OtpGenServer: like gen\_server module

# Extensions



V1.1

V1.0

V2.0

# Extensions

## .near

```
Special value type:  
#geo(lat :: float(), lon :: float())  
  
9> lucene:add(  
    [{"index", 1},{location, lucene_utils:geo(1/2,1+1.5)}]  
    || I <- lists:seq(-10,10)).  
ok  
parameters: lat,lon,dist (in miles)  
  
10> lucene:match("location.near:0.0,1.0,100", 3).  
{{{"index",0},  
  [{"location,{geo,-8.381903171539307e-8,1.5000000782310963}],  
   {'score',-17.292743682861328}],  
  {"index",-1},  
  [{"location,{geo,-0.5000000540167093,0.500000137835741}],  
   {'score',-24.45547866821289}],  
  {"total_hits",2},  
  {"first_hit",1},  
  {"query_time",1984},  
  {"search_time",1021}}}
```

## .erlang

```
arguments (last one is always a list  
of values for the chosen field)  
an Erlang function  
  
11> lucene:match(  
    "index.erlang:\\"lists:map:[fun(X) -> X*1.0 end]\\"", 3).  
{{{"index",10},  
  [{"location,{geo,4.99999953433871,11.500000152736902}},{  
   {'score',5.0}],  
  {"index",9},  
  [{"location,{geo,4.49999983236194,10.49999987706542}},{  
   {'score',4.499999523162842}],  
  {"index",8},  
  [{"location,{geo,4.000000013038516,9.499999936670065}},{  
   {'score',3.999999761581421}],  
  {"total_hits",21},  
  {"first_hit",1},  
  {"query_time",2256},  
  {"search_time",1240},  
  {next_page,<<172,237,0,5,115,114,0,36,99,111,109,46,  
   116,105,103,101,114,116,101,120,...>>]}}
```

# .near

Special value type:  
#geo(lat :: float(), lon :: float())

```
9> lucene:add(  
    [[{index, I},{location, lucene_utils:geo(I/2,I+1.5)}]  
     || I <- lists:seq(-10,10)]).  
ok  
parameters: lat,lon,dist (in miles)  
  
10> lucene:match("location.near:0.0,1.0,100", 3).  
[[{{index,0},  
    {location,{geo,-8.381903171539307e-8,1.5000000782310963}},  
    {'score',-17.292743682861328}],  
 [{index,-1},  
  {location,{geo,-0.5000000540167093,0.500000137835741}},  
  {'score',-24.45547866821289}],  
 [{total_hits,2},  
  {first_hit,1},  
  {query_time,1984},  
  {search_time,1021}]]
```

# .erlang

an Erlang function

arguments (last one is always a list  
of values for the chosen field)

```
11> lucene:match(  
    "index.erlang:\\"lists:map:[fun(X) -> X*1.0 end]\\"", 3).  
[[[{{index,10},  
     {location,{geo,4.99999953433871,11.5},  
      {'score',5.0}},  
    {{index,9},  
     {location,{geo,4.49999983236194,10.4999987706542}},  
      {'score',4.49999523162842}},  
    {{index,8},  
     {location,{geo,4.0,9.49999936670065}},  
      {'score',3.99999761581421}}],  
    [{total_hits,21},  
     {first_hit,1},  
     {query_time,2256},  
     {search_time,1240},  
     {next_page,<<172,237,0,5,115,114,0,36,99,111,109,46,  
      116,105,103,101,114,116,101,120,...>>}]]
```

# Extensions

## Extensions

```
public LuceneServer(OtpNode host, int allowedThreads)
    throws CorruptIndexException, LockObtainFailedException,
    IOException {
    super(host, "lucene_server");
    ...
    Extensions ext = new Extensions("ext");
    ext.add("near", new NearParserExtension());
    ext.add("erlang", new ErlangParserExtension(this.translator));
    this.extensions = ext;
    ...
}

private QueryParser queryParser() {
    return new LuceneQueryParser(Version.LUCENE_36, this.analyzer,
        this.translator, this.extensions);
}
```

is used when building query parser

### .near

- NearParserExtension is built around lucene-spatial library
- Indexes fields using n-tiers
- We added a parser for #geo{} records

### .erlang



# Extensions

```
public LuceneServer(OtpNode host, int allowedThreads)
    throws CorruptIndexException, LockObtainFailedException,
           IOException {
    super(host, "lucene_server");
    ...
    Extensions ext = new Extensions('.');
    ext.add("near", new NearParserExtension());
    ext.add("erlang", new ErlangParserExtension(this.translator));
    this.extensions = ext;
    ...
}

private QueryParser queryParser() {
    return new LuceneQueryParser(Version.LUCENE_36, this.analyzer,
        this.translator, this.extensions);
}
```

an extensions array  
using '.' as delimiter

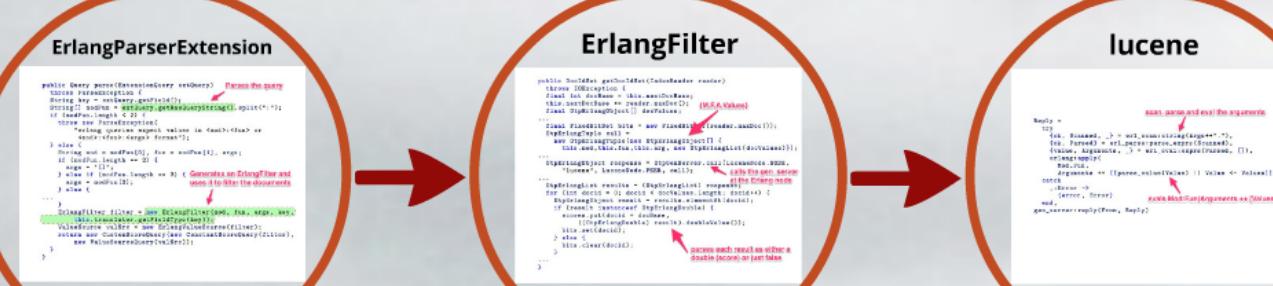
is used when building  
query parser

# .near

- NearParserExtension is built around lucene-spatial library
- Indexes fields using n-tiers
- We added a parser for #geo{} records



# .erlang



# ErlangParserExtension

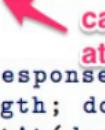
```
public Query parse(ExtensionQuery extQuery)          Parses the query
throws ParseException {
String key = extQuery.getField();
String[] modFun = extQuery.getRawQueryString().split(":");
if (modFun.length < 2) {
    throw new ParseException(
        "erlang queries expect values in <mod>:<fun> or
        <mod>:<fun>:<args> format");
} else {
    String mod = modFun[0], fun = modFun[1], args;
    if (modFun.length == 2) {
        args = "[]";
    } else if (modFun.length == 3) { Generates an ErlangFilter and
        args = modFun[2];
    } else {
        ...
    }
    ErlangFilter filter = new ErlangFilter(mod, fun, args, key,
        this.translator.getFieldType(key));
    ValueSource valSrc = new ErlangValueSource(filter);
    return new CustomScoreQuery(new ConstantScoreQuery(filter),
        new ValueSourceQuery(valSrc));
}
}
```



# ErlangFilter

```
public DocIdSet getDocIdSet(IndexReader reader)
throws IOException {
final int docBase = this.nextDocBase;
this.nextDocBase += reader.maxDoc();
final OtpErlangObject[] docValues;
...
final FixedBitSet bits = new FixedBitSet(reader.maxDoc());
OtpErlangTuple call =
new OtpErlangTuple(new OtpErlangObject[] {
    this.mod, this.fun, this.arg, new OtpErlangList(docValues)});
...
OtpErlangObject response = OtpGenServer.call(LuceneNode.NODE,
    "lucene", LuceneNode.PEER, call);
...
OtpErlangList results = (OtpErlangList) response;
for (int docid = 0; docid < docValues.length; docid++) {
    OtpErlangObject result = results.elementAt(docid);
    if (result instanceof OtpErlangDouble) {
        scores.put(docid + docBase,
            ((OtpErlangDouble) result).doubleValue());
        bits.set(docid);
    } else {
        bits.clear(docid);
    }
...
}
```

{M,F,A,Values} 

calls the gen\_server at the Erlang node 

parses each result as either a double (score) or just false 

# lucene

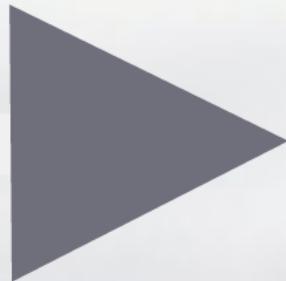


```
scan, parse and eval the arguments
Reply =
try
  {ok, Scanned, _} = erl_scan:string(Args++"."),
  {ok, Parsed} = erl_parse:parse_exprs(Scanned),
  {value, Arguments, _} = erl_eval:exprs(Parsed, []),
  erlang:apply(
    Mod, Fun,
    Arguments ++ [[parse_value(Value) || Value <- Values]])
catch
  _:Error ->
  {error, Error}
end,
gen_server:reply(From, Reply)
```

evals Mod:Fun(Arguments ++ [Values])

1

V2.0



## The Future

- More extensions!
- Security improvements
- Other backends? (Currently it's just in-memory)
- More users! That's YOU!!

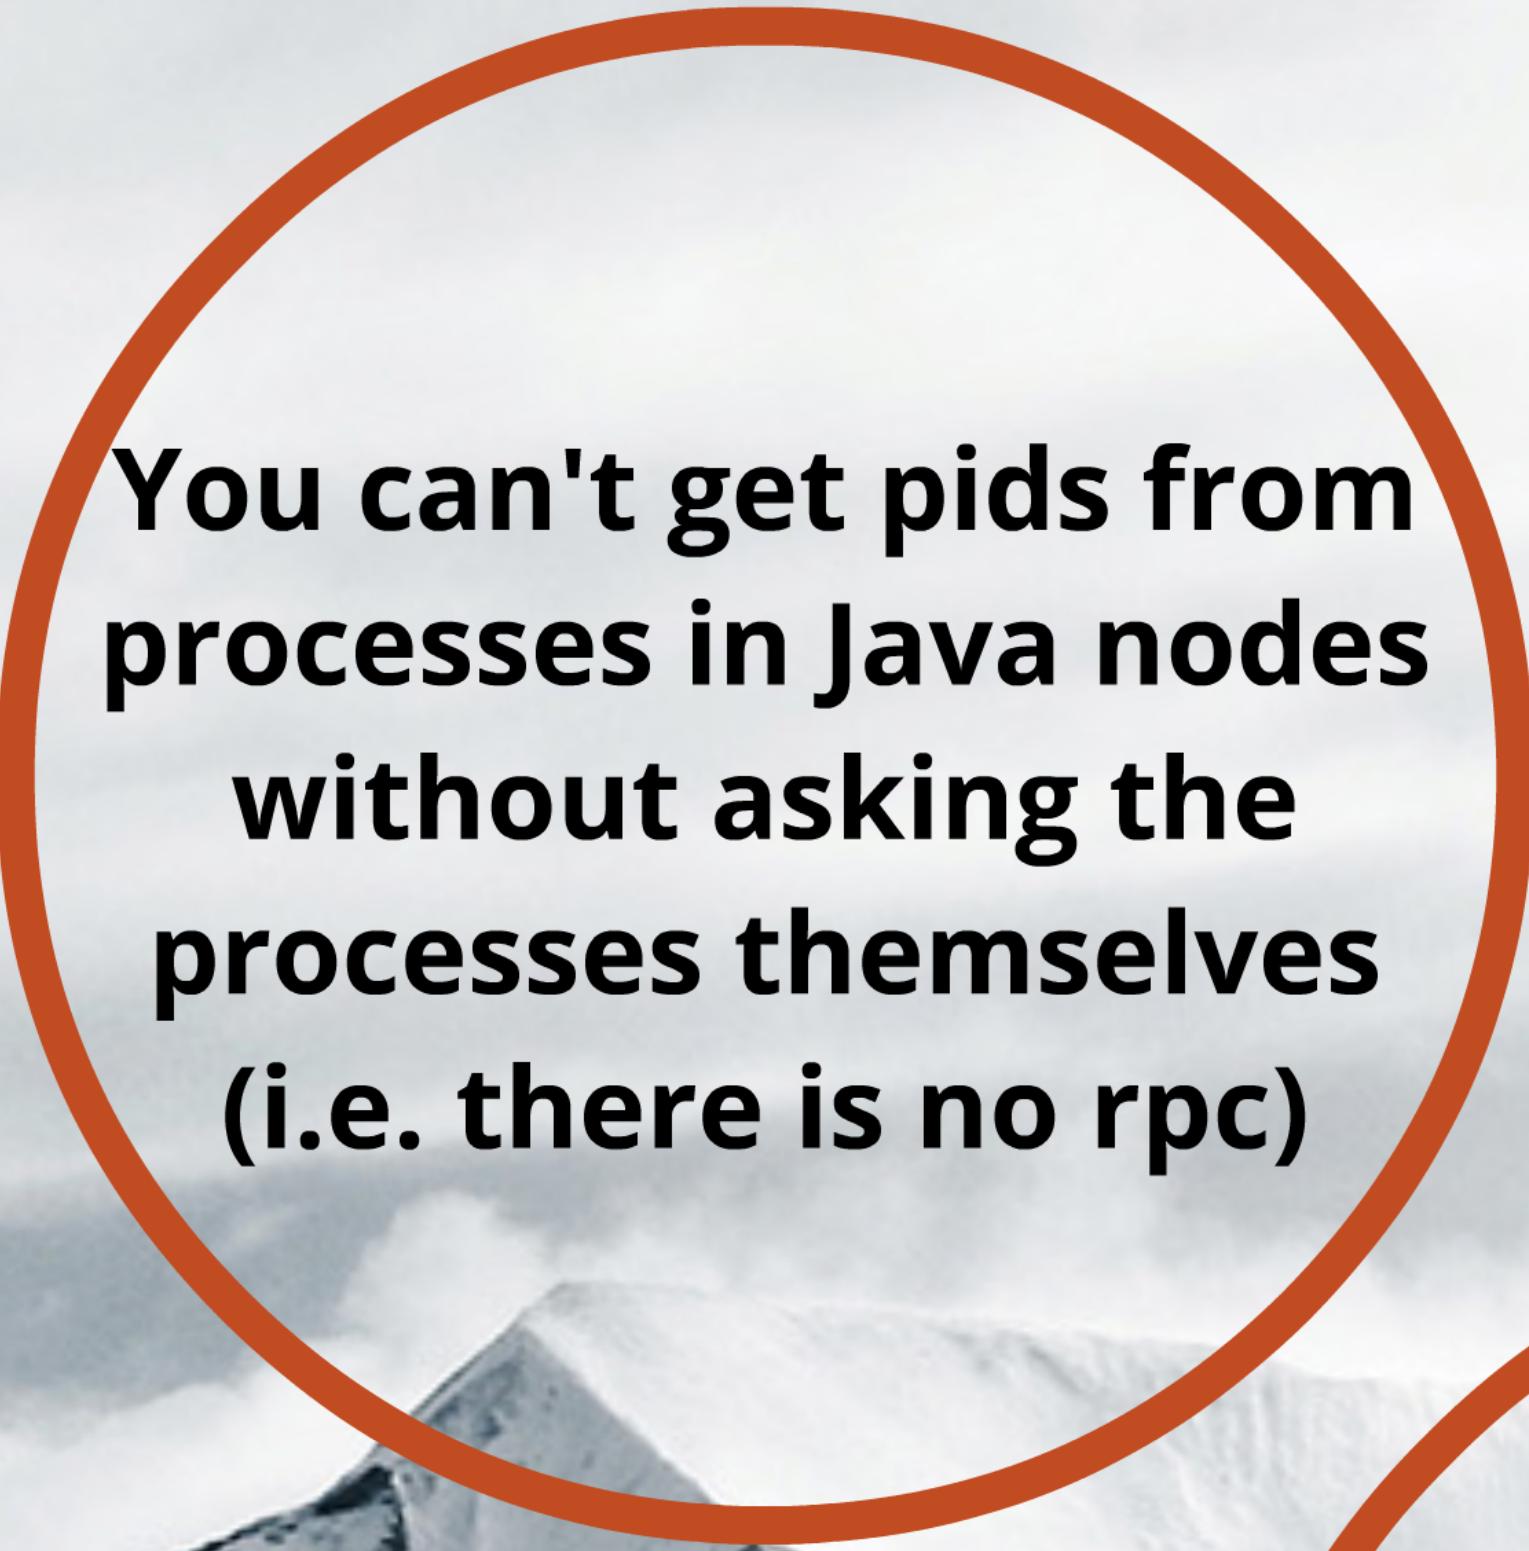


## *Lessons Learned*

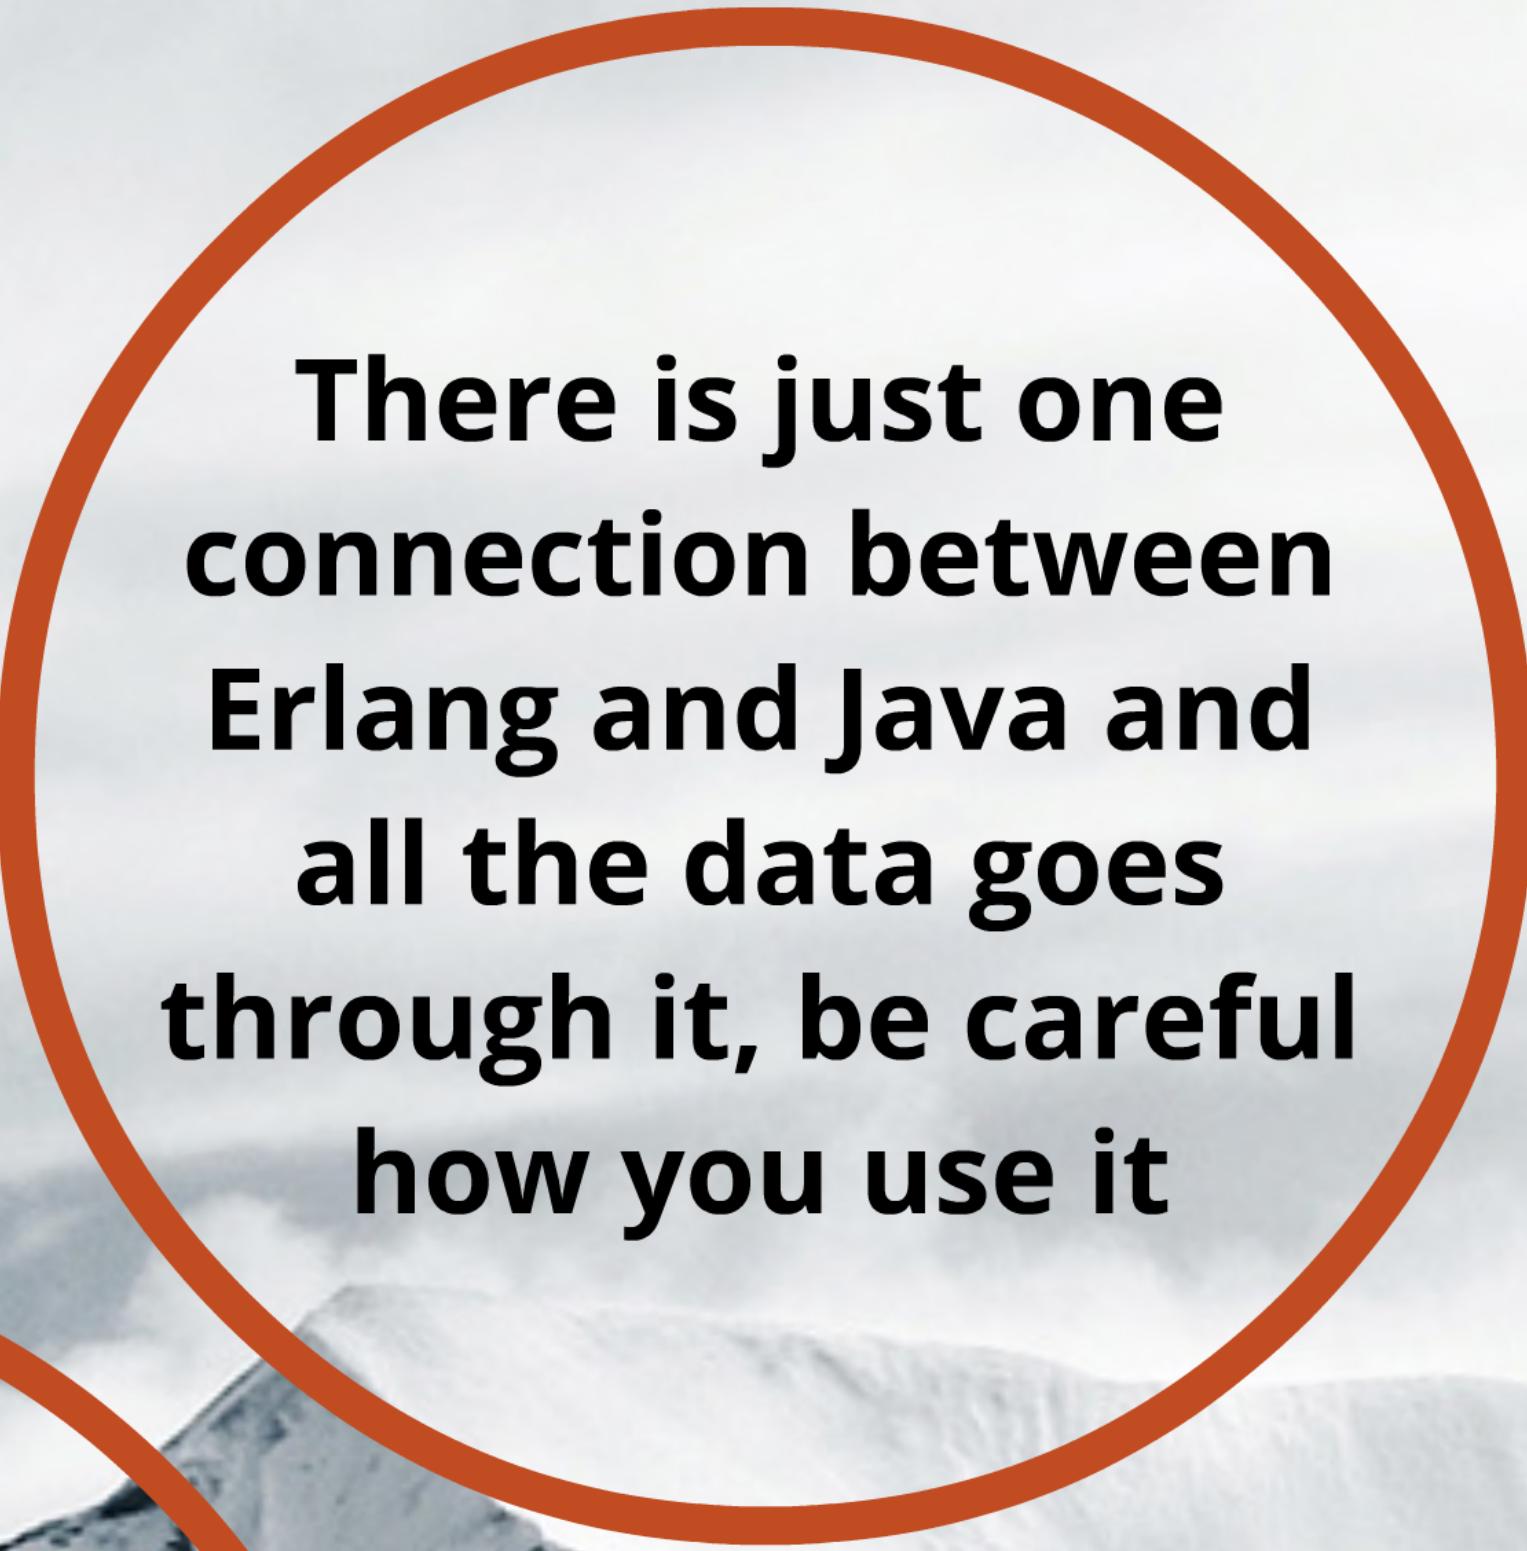
You can't get pids from processes in Java nodes without asking the processes themselves (i.e. there is no rpc)

There is just one connection between Erlang and Java and all the data goes through it, be careful how you use it

Not every string in Erlang is an OtpErlangString in Java. Lists longer than 65535 elements are encoded as OtpErlangList.



**You can't get pids from  
processes in Java nodes  
without asking the  
processes themselves  
(i.e. there is no rpc)**



**There is just one  
connection between  
Erlang and Java and  
all the data goes  
through it, be careful  
how you use it**



**Not every string in  
Erlang is an  
OtpErlangString in Java.  
Lists longer than 65535  
elements are encoded  
as OtpErlangList.**

# *Resources*

- [github.com/tigertext/lucene\\_server](https://github.com/tigertext/lucene_server)
- [github.com/inaka/jinterface\\_stdlib](https://github.com/inaka/jinterface_stdlib)
- [lucene.apache.org/core](http://lucene.apache.org/core)
- [about.me/elbrujohalcon](http://about.me/elbrujohalcon)
- [github.com/inaka](https://github.com/inaka)
- [inaka.net](http://inaka.net)
- [inaka.net/blog](http://inaka.net/blog)



*THANK YOU VERY MUCH!!*